


Li Yonghao (Assignment 4)

1. Load the data and visualization

Amount of images 6929: Artifacts: 2380, cancer: 2878, normal: 1187, other:484.

I used to use keras, so I use the `os.listdir()` and `img_to_array()` function to load the data, then use the `np.transpose()` to reshape the data fit for pytorch.



`imshow()` to visualize the data

I add labels to each category, artifacts: 0, cancer: 1, normal: 2, other:3.

Then I extract 100 images of each category as test set, rest will be training set and validation set with the rate of 0.8.

I normalize the data in to [0,1] to get a better result of training.

Below is my CNN model with 3 conv2d layers and 2 linear layers:

```
class cnn(torch.nn.Module):
    def __init__(self):
        super().__init__()

        self.cnn_layers = torch.nn.Sequential(

            # Defining a 2D convolution layer
            torch.nn.Conv2d(3, 64, kernel_size=3, stride=2, padding=1),
            torch.nn.BatchNorm2d(64),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2, stride=2),

            # Defining another 2D convolution layer
            torch.nn.Conv2d(64, 128, kernel_size=3, stride=2, padding=1),
            torch.nn.BatchNorm2d(128),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2, stride=2),

            torch.nn.Conv2d(128, 256, kernel_size=3, stride=2, padding=1),
            torch.nn.BatchNorm2d(256),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2, stride=2)

        )

        self.linear_layers = torch.nn.Sequential(
            torch.nn.Flatten(1),
            torch.nn.Linear(256 * 2 * 2, 512),
            torch.nn.ReLU(inplace=True),
            torch.nn.Dropout(p=0.2),
            torch.nn.Linear(512, 4)
        )

    # Defining the forward pass
    def forward(self, x):
        x = self.cnn_layers(x)
        x = x.view(x.size(0), -1)
        x = self.linear_layers(x)
        return x
```

For training, I try different parameters:

Learning rate: 0.1, 0.01, 0.001 (best one with lowest loss score under 0.1).

Rate decay: I use Adam and try Adagrad, Adadelata, RMSprop and Time-Based Decay: $lr *= (1. / (1. + self.decay * self.iterations))$ in SGD optimizer.

Drop out: I try 0.1, 0.2 (best one), 0.3, 0.4, 0.5.

Weight decay: 0.01


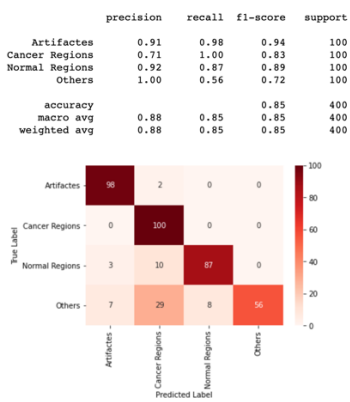
Early stop: I do then by using loss of validation and training, if the validation loss doesn't decrease but increase, I will save the best model to avoid that.

Network architecture adjustment: I try some and use the best one of my tries.

Detail of training:

```
import torch.optim as optim
loss_func = torch.nn.CrossEntropyLoss()
lr = 0.001
weight_decay=0.01
optimizer = optim.Adam(net.parameters(), lr=lr, betas=(0.9, 0.999), eps=1e-08, weight_decay=weight_decay)
optimizer_name = 'Adam'
```

```
epoch5 train loss: 0.10503230807257861, val loss: 0.1757908747954802 val accuracy: 0.947166921898928, patience: 2
epoch6 train loss: 0.09976871511558207, val loss: 0.1662780845707113 val accuracy: 0.9418070444104135, patience: 3
epoch7 train loss: 0.0749298469140762, val loss: 0.14441461962732402 val accuracy: 0.9448698315467075, patience: 4
epoch8 train loss: 0.06901015564617587, val loss: 0.17205342111934314 val accuracy: 0.9402756508422665, patience: 5
epoch9 train loss: 0.09813730950217421, val loss: 0.1789132058620453 val accuracy: 0.947166921898928, patience: 6
epoch10 train loss: 0.11122148256839776, val loss: 0.17234653302214362 val accuracy: 0.9464012251148545, patience: 7
epoch11 train loss: 0.0806708085555129, val loss: 0.1486456888643178 val accuracy: 0.9433384379785605, patience: 8
epoch12 train loss: 0.0775321738159422, val loss: 0.16138073801994324 val accuracy: 0.9486983154670751, patience: 9
epoch13 train loss: 0.08681874749500577, val loss: 0.1253729723393917 val accuracy: 0.955895865237366, patience: 0
epoch14 train loss: 0.1136283047264959, val loss: 0.1749492881679535 val accuracy: 0.9433384379785605, patience: 1
epoch15 train loss: 0.07273438318473537, val loss: 0.146782631223852 val accuracy: 0.9509954058192955, patience: 2
epoch16 train loss: 0.06373720738764216, val loss: 0.15071229162541303 val accuracy: 0.9532924961715161, patience: 3
epoch17 train loss: 0.0997676791444057, val loss: 0.1840338801795786 val accuracy: 0.9395995405819295, patience: 4
epoch18 train loss: 0.08295710430276103, val loss: 0.15275009721517563 val accuracy: 0.9448698315467075, patience: 5
epoch19 train loss: 0.05835777456756627, val loss: 0.1428982056333802 val accuracy: 0.9464012251148545, patience: 6
epoch20 train loss: 0.04733092325628425, val loss: 0.162048642155257 val accuracy: 0.9486983154670751, patience: 7
epoch21 train loss: 0.05494103239985501, val loss: 0.15560766444946844 val accuracy: 0.9509954058192955, patience: 8
epoch22 train loss: 0.1178471514182467, val loss: 0.240866535089227 val accuracy: 0.9241960183767228, patience: 9
epoch23 train loss: 0.06839014062794244, val loss: 0.2021694074977528 val accuracy: 0.9249617151607963, patience: 10
epoch24 train loss: 0.06627383128535456, val loss: 0.21138350665569305 val accuracy: 0.9234303215926493, patience: 11
epoch25 train loss: 0.0741499169402355, val loss: 0.16271070323207162 val accuracy: 0.9479326186830015, patience: 12
epoch26 train loss: 0.07667812802224624, val loss: 0.14247851276939564 val accuracy: 0.9479326186830015, patience: 13
epoch27 train loss: 0.07040098373119424, val loss: 0.19290945340286603 val accuracy: 0.9379785604900459, patience: 14
epoch28 train loss: 0.0468915681176275, val loss: 0.12374418804591353 val accuracy: 0.9494640122511485, patience: 15
epoch29 train loss: 0.042985512034558665, val loss: 0.18533993580124594 val accuracy: 0.9479326186830015, patience: 16
epoch30 train loss: 0.06011595436167426, val loss: 0.12141064757650549 val accuracy: 0.9609494640122511, patience: 0
epoch31 train loss: 0.04567546397447586, val loss: 0.1438532363284718 val accuracy: 0.955895865237366, patience: 1
epoch32 train loss: 0.04597299650493192, val loss: 0.17830072275616907 val accuracy: 0.931852862174579, patience: 2
epoch33 train loss: 0.0494037525864636, val loss: 0.19506299766627225 val accuracy: 0.9226464248085758, patience: 3
epoch34 train loss: 0.058943349214676, val loss: 0.19341136796070737 val accuracy: 0.9379785604900459, patience: 4
epoch35 train loss: 0.07919367321017312, val loss: 0.15581736781380393 val accuracy: 0.9464012251148545, patience: 5
epoch36 train loss: 0.04447611475863108, val loss: 0.16138289462436328 val accuracy: 0.945635528330781, patience: 6
epoch37 train loss: 0.0340266479624806, val loss: 0.1312564733353528 val accuracy: 0.9548238897396631, patience: 7
epoch38 train loss: 0.049163964009139596, val loss: 0.21083865653384815 val accuracy: 0.9133843797856048, patience: 8
epoch39 train loss: 0.0660683104722965, val loss: 0.15201217274774204 val accuracy: 0.9532924961715161, patience: 9
```



When I reload the val path of best model, I get the test accuracy of 95%. But if just use the final model after all epochs it will be 85%. (Early stop make sense)

Test accuracy vs size of train set without early stop: (less image can get high train accuracy, but test accuracy very low)

